

```
In [54]: # 1.Importing the Libraries
import sys
from ipykernel import kernelapp as app
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from numpy import percentile
from numpy.random import rand
from scipy import stats
from sklearn.metrics import *
import seaborn as sn
```

```
In [26]: # 2.Read the data as a data frame
dataset = pd.read_csv('E:/IISWBM_Materials/2nd_sem/Advance_analytics/bank-full.csv')

bankdata = pd.DataFrame(dataset)

print(bankdata)
```

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	
...	...	...	...	...	...	...	...	...	
45206	51	technician	married	tertiary	no	825	no	no	
45207	71	retired	divorced	primary	no	1729	no	no	
45208	72	retired	married	secondary	no	5715	no	no	
45209	57	blue-collar	married	secondary	no	668	no	no	
45210	37	entrepreneur	married	secondary	no	2971	no	no	

	contact	day	month	duration	campaign	pdays	previous	poutcome	\
0	unknown	5	may	261	1	-1	0	unknown	
1	unknown	5	may	151	1	-1	0	unknown	
2	unknown	5	may	76	1	-1	0	unknown	
3	unknown	5	may	92	1	-1	0	unknown	
4	unknown	5	may	198	1	-1	0	unknown	
...	...	...	...	...	...	...	...	...	
45206	cellular	17	nov	977	3	-1	0	unknown	
45207	cellular	17	nov	456	2	-1	0	unknown	
45208	cellular	17	nov	1127	5	184	3	success	
45209	telephone	17	nov	508	4	-1	0	unknown	
45210	cellular	17	nov	361	2	188	11	other	

	Target
0	no
1	no
2	no
3	no
4	no
...	...
45206	yes
45207	yes
45208	yes
45209	no
45210	no

[45211 rows x 17 columns]

```
In [55]: # 3.a.Shape of the data

shape = bankdata.shape
print(shape)

(45211, 17)
```

In [56]: *# 3.b.Data type of each attribute*

```
datatypes = bankdata.dtypes
print(datatypes)
```

```
age          int64
job          object
marital      object
education    object
default      object
balance      int64
housing      object
loan         object
contact      object
day          int64
month        object
duration     int64
campaign     int64
pdays       int64
previous     int64
poutcome     object
Target       object
dtype: object
```

In [57]: *# 3.c.Checking the presence of missing values and get rid of those missing values*

```
bankdata.isnull()
bankdata.dropna()
bankdata = bankdata.fillna(100)
int_df = bankdata.select_dtypes(include=['int64']).copy()
float_df=bankdata.select_dtypes(include=['float64']).copy()
bankdata_int_float = pd.concat([float_df,int_df], axis=1, join_axes=[int_df.index])
obj_df = bankdata.select_dtypes(include=['object']).copy()
obj_df.head()
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
le.fit(obj_df["job"].astype(str))
list(le.classes_)
obj_df_trf=obj_df.astype(str).apply(le.fit_transform)
bankdata_final = pd.concat([bankdata_int_float,obj_df_trf], axis=1, join_axes=[bankdata_int_float.index])
bankdata_final.head()
X = bankdata_final.iloc[:,0:16].values
print(X)
y = bankdata_final.iloc[:, 16].values
print(y)
```

D:\Conda\Anaconda3\lib\site-packages\ipykernel\_launcher.py:7: FutureWarning: The join\_axes-keyword is deprecated. Use .reindex or .reindex\_like on the result to achieve the same functionality.

```
import sys
```

D:\Conda\Anaconda3\lib\site-packages\ipykernel\_launcher.py:15: FutureWarning: The join\_axes-keyword is deprecated. Use .reindex or .reindex\_like on the result to achieve the same functionality.

```
from ipykernel import kernelapp as app
```

```
[[ 58 2143    5 ...    2    8    3]
 [ 44   29    5 ...    2    8    3]
 [ 33    2    5 ...    2    8    3]
 ...
 [ 72 5715   17 ...    0    9    2]
 [ 57  668   17 ...    1    9    3]
 [ 37 2971   17 ...    0    9    1]]
[0 0 0 ... 1 0 0]
```

```
In [58]: # 3.d 5point summary of numerical attributes
age = bankdata.iloc[:,0].values
balance = bankdata.iloc[:,5].values
day = bankdata.iloc[:,9].values
duration = bankdata.iloc[:,11].values
campaign = bankdata.iloc[:,12].values
Pday = bankdata.iloc[:,13].values
age_quartiles = percentile(age, [25, 50, 75])
age_min, age_max = age.min(), age.max()
bal_quartiles = percentile(balance, [25, 50, 75])
bal_min, bal_max = balance.min(), balance.max()
day_quartiles = percentile(day, [25, 50, 75])
day_min, day_max = day.min(), day.max()
duration_quartiles = percentile(duration, [25, 50, 75])
duration_min, duration_max = duration.min(), duration.max()
campaign_quartiles = percentile(campaign, [25, 50, 75])
campaign_min, campaign_max = campaign.min(), campaign.max()
Pday_quartiles = percentile(Pday, [25, 50, 75])
Pday_min, Pday_max = Pday.min(), Pday.max()
print('AGE: ', 'Min: %.3f', age_min, 'Q1: %.3f', age_quartiles[0], 'Median: %.3f', age_quartiles[1],
      'Q3: %.3f', age_quartiles[2], 'Max: %.3f', age_max)
print('BALANCE: ', 'Min: %.3f', bal_min, 'Q1: %.3f', bal_quartiles[0], 'Median: %.3f', bal_quartiles[1],
      'Q3: %.3f', bal_quartiles[2], 'Max: %.3f', bal_max)
print('DAY: ', 'Min: %.3f', day_min, 'Q1: %.3f', day_quartiles[0], 'Median: %.3f', day_quartiles[1],
      'Q3: %.3f', day_quartiles[2], 'Max: %.3f', day_max)
print('DURATION: ', 'Min: %.3f', duration_min, 'Q1: %.3f', duration_quartiles[0], 'Median: %.3f', duration_quartiles[1],
      'Q3: %.3f', duration_quartiles[2], 'Max: %.3f', duration_max)
print('CAMPAIGN: ', 'Min: %.3f', campaign_min, 'Q1: %.3f', campaign_quartiles[0], 'Median: %.3f', campaign_quartiles[1],
      'Q3: %.3f', campaign_quartiles[2], 'Max: %.3f', campaign_max)
print('PDAY: ', 'Min: %.3f', Pday_min, 'Q1: %.3f', Pday_quartiles[0], 'Median: %.3f', Pday_quartiles[1],
      'Q3: %.3f', Pday_quartiles[2], 'Max: %.3f', Pday_max)
```

```
AGE:  Min: %.3f 18 Q1: %.3f 33.0 Median: %.3f 39.0 Q3: %.3f 48.0 Max: %.3f 95
BALANCE:  Min: %.3f -8019 Q1: %.3f 72.0 Median: %.3f 448.0 Q3: %.3f 1428.0 Max: %.3f 102127
DAY:  Min: %.3f 1 Q1: %.3f 8.0 Median: %.3f 16.0 Q3: %.3f 21.0 Max: %.3f 31
DURATION:  Min: %.3f 0 Q1: %.3f 103.0 Median: %.3f 180.0 Q3: %.3f 319.0 Max: %.3f 4918
CAMPAIGN:  Min: %.3f 1 Q1: %.3f 1.0 Median: %.3f 2.0 Q3: %.3f 3.0 Max: %.3f 63
PDAY:  Min: %.3f -1 Q1: %.3f -1.0 Median: %.3f -1.0 Q3: %.3f -1.0 Max: %.3f 871
```

```
In [59]: # 3.e. Checking the presence of outliers
def drop_numerical_outliers(bankdata_final, z_thresh=3):
    # Constrains will contain `True` or `False` depending on if it is a value below the threshold.
    constrains = bankdata_final.select_dtypes(include=[np.number]).apply(lambda x: np.abs(stats.zscore(x)) < z_thresh,
                                reduce=False).all(axis=1)
    # Drop (inplace) values set to be rejected
    bankdata_final.drop(bankdata_final.index[~constrains], inplace=True)
```

```
In [60]: dropdata = drop_numerical_outliers(bankdata_final)
print(dropdata)
```

None

D:\Conda\Anaconda3\lib\site-packages\ipykernel\_launcher.py:4: FutureWarning: The reduce argument is deprecated and will be removed in a future version. You can specify result\_type='reduce' to try to reduce the result to the original dimensions after removing the cwd from sys.path.

```
In [61]: # 4. Prepare the data to train a model - check if data types are appropriate, get rid of the missing values etc
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.35, random_state = 0)

from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
k_fold = KFold(n_splits=10, shuffle=True, random_state=0)

#X_train.head()
from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)
```


```
In [62]: # Fitting Random Forest Classification to the Training set
from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(n_estimators = 10, criterion = 'entropy', random_state = 0)
classifier.fit(X_train, y_train)

# Predicting the Test set results
y_pred = classifier.predict(X_test)
predrfc = classifier.predict_proba(X_test)
# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print(cm)
Acc = accuracy_score(y_test, y_pred)
print("Accuracy score of Random Forest::: ",Acc)
cmrfc = pd.DataFrame(cm, columns=np.unique(y_test), index = np.unique(y_test))
cmrfc.index.name = 'Actual'
cmrfc.columns.name = 'Predicted'
plt.figure(figsize = (10,7))
sn.set(font_scale=1.4)#for label size
sn.heatmap(cm, annot=True)

plt.show()
RFCV = (cross_val_score(classifier, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
```

```
[[13503  458]
 [ 1210  653]]
```

Accuracy score of Random Forest::: 0.8945904954499494




```
In [63]: # Fitting Decision Tree Classification to the Training set
from sklearn.tree import DecisionTreeClassifier
classifier1 = DecisionTreeClassifier(criterion = 'entropy', random_state = 0)
classifier1.fit(X_train, y_train)

# Predicting the Test set results
y_pred1 = classifier1.predict(X_test)
predproba = classifier1.predict_proba(X_test)
# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm1 = confusion_matrix(y_test, y_pred1)
print(cm1)
Acc = accuracy_score(y_test, y_pred1)
print("Accuracy score of DecisionTree::: ",Acc)
cmdtr = pd.DataFrame(cm1, columns=np.unique(y_test), index = np.unique(y_test))
cmdtr.index.name = 'Actual'
cmdtr.columns.name = 'Predicted'
plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(cm1, annot=True)

plt.show()
DTREECV = (cross_val_score(classifier1, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
```

```
[[12965  996]
 [ 1006  857]]
```

Accuracy score of DecisionTree::: 0.8734833164812943




```
In [64]: # Fitting K-NN to the Training set
from sklearn.neighbors import KNeighborsClassifier
classifier2 = KNeighborsClassifier(n_neighbors = 5, metric = 'minkowski', p = 2)
classifier2.fit(X_train, y_train)

# Predicting the Test set results
y_pred2 = classifier.predict(X_test)
predknn = classifier.predict_proba(X_test)

# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm2 = confusion_matrix(y_test, y_pred2)
print(cm2)
Acc = accuracy_score(y_test, y_pred2)
print("Accuracy score of KNN::: ",Acc)
cmknn = pd.DataFrame(cm2, columns=np.unique(y_test), index = np.unique(y_test))
cmknn.index.name = 'Actual'
cmknn.columns.name = 'Predicted'
plt.figure(figsize = (10,7))
sn.set(font_scale=1.4)#for label size
sn.heatmap(cm2, annot=True)

plt.show()
KNNCV = (cross_val_score(classifier2, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
```

```
[[13503  458]
 [ 1210  653]]
Accuracy score of KNN:::  0.8945904954499494
```



```
In [65]: from sklearn.linear_model import LogisticRegression
logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)
logpred = logmodel.predict(X_test)

cm3 = confusion_matrix(y_test, logpred)
print(cm3)
Acc = accuracy_score(y_test, logpred)
print("Accuracy score of Logistic::: ",Acc)
cmlog = pd.DataFrame(cm2, columns=np.unique(y_test), index = np.unique(y_test))
cmlog.index.name = 'Actual'
cmlog.columns.name = 'Predicted'
plt.figure(figsize = (10,7))
sn.set(font_scale=1.4)#for label size
sn.heatmap(cm3, annot=True)

plt.show()
LOGCV = (cross_val_score(logmodel, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
```


D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

[[13654 307]

[ 1453 410]]

Accuracy score of Logistic::: 0.8887765419615774



D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.


FutureWarning)

```
In [66]: # Support vector machine Model
from sklearn.svm import SVC
svc= SVC(kernel = 'sigmoid')
svc.fit(X_train, y_train)
svcpred = svc.predict(X_test)
cm4 = confusion_matrix(y_test, svcpred)
print(cm4)
acc4 = round(accuracy_score(y_test, svcpred),2)*100
print(acc4)
#SVCCV = (cross_val_score(svc, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
cmsvc = pd.DataFrame(cm4, columns=np.unique(y_test), index = np.unique(y_test))
cmsvc.index.name = 'Actual'
cmsvc.columns.name = 'Predicted'
plt.figure(figsize = (10,7))
sn.set(font_scale=1.4)#for label size
sn.heatmap(cm4, annot=True)

plt.show()
SVCCV = (cross_val_score(svc, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
```



```
[[12720 1241]
 [ 1354  509]]
84.0
```



D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)


D:\Conda\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

```
In [67]: # Naive Bayes Classifier Model
from sklearn.naive_bayes import GaussianNB
gaussiannb= GaussianNB()
gaussiannb.fit(X_train, y_train)
gaussiannbpred = gaussiannb.predict(X_test)
probs = gaussiannb.predict(X_test)
cm5 = confusion_matrix(y_test, probs)
print(cm5)
acc5 = round(accuracy_score(y_test, probs),2)*100
print(acc5)
#SVCCV = (cross_val_score(svc, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
cmnv = pd.DataFrame(cm5, columns=np.unique(y_test), index = np.unique(y_test))
cmnv.index.name = 'Actual'
cmnv.columns.name = 'Predicted'
plt.figure(figsize = (10,7))
sn.set(font_scale=1.4)#for label size
sn.heatmap(cm5, annot=True)

plt.show()
GAUSIAN = (cross_val_score(gaussiannb, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
```

```
[[12386  1575]
 [   983   880]]
84.0
```

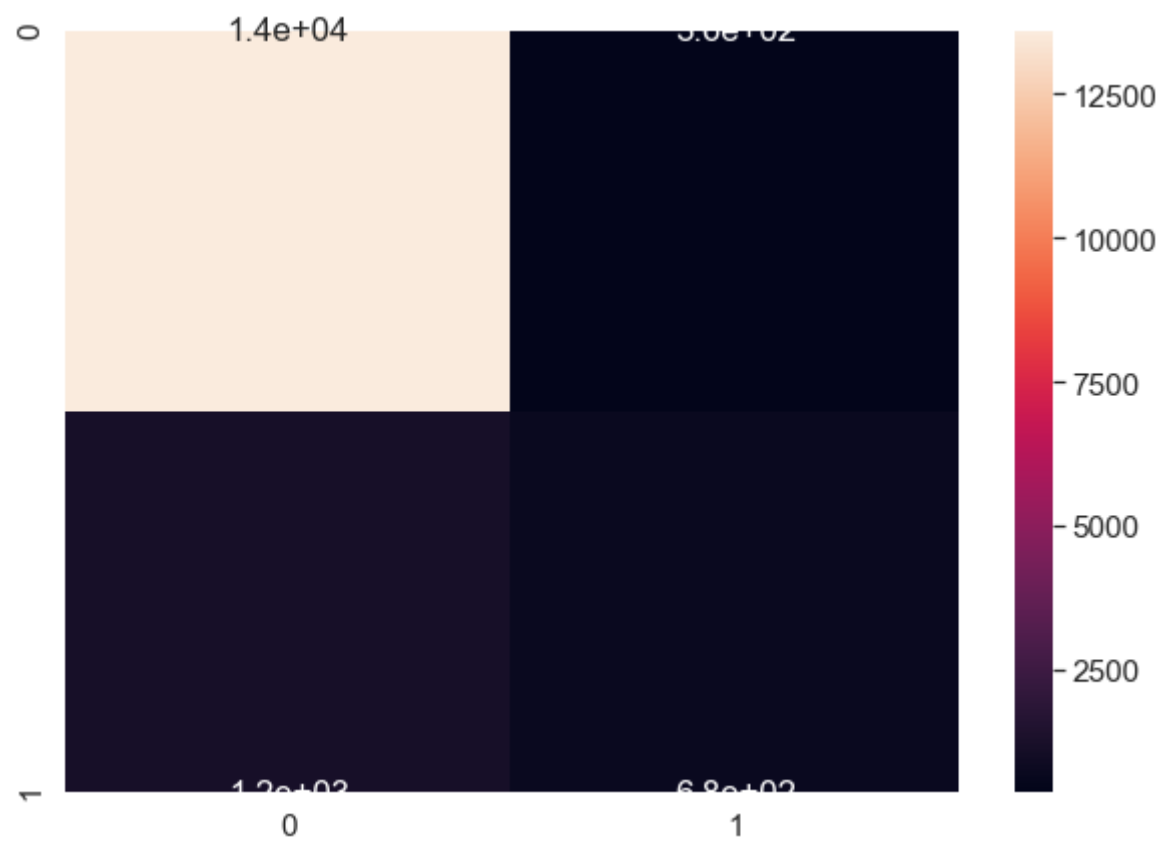


```
In [68]: # XGBOOST Classifier Model
from xgboost import XGBClassifier
xgb = XGBClassifier()
xgb.fit(X_train, y_train)
xgbprd = xgb.predict(X_test)

cm6 = confusion_matrix(y_test, xgbprd)
print(cm6)
acc6 = round(accuracy_score(y_test, xgbprd),2)*100
print(acc6)
#SVCCV = (cross_val_score(svc, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
cmxg = pd.DataFrame(cm6, columns=np.unique(y_test), index = np.unique(y_test))
cmxg.index.name = 'Actual'
cmxg.columns.name = 'Predicted'
plt.figure(figsize = (10,7))
sn.set(font_scale=1.4)#for label size
sn.heatmap(cm6, annot=True)

plt.show()
XGB = (cross_val_score(estimator = xgb, X = X_train, y = y_train, cv = 10).mean())
```

```
[[13596  365]
 [ 1184  679]]
90.0
```




```
In [69]: # Gradient Boosting Model

from sklearn.ensemble import GradientBoostingClassifier
gbk = GradientBoostingClassifier()
gbk.fit(X_train, y_train)
gbkpred = gbk.predict(X_test)

cm7 = confusion_matrix(y_test, gbkpred)
print(cm7)
acc7 = round(accuracy_score(y_test, gbkpred),2)*100
print(acc7)
#SVCCV = (cross_val_score(svc, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
cmgbk = pd.DataFrame(cm7, columns=np.unique(y_test), index = np.unique(y_test))
cmgbk.index.name = 'Actual'
cmgbk.columns.name = 'Predicted'
plt.figure(figsize = (10,7))
sn.set(font_scale=1.4)#for label size
sn.heatmap(cm7, annot=True)

plt.show()
GBKCV = (cross_val_score(gbk, X_train, y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy').mean())
```



```
In [70]: # Comparison of Different Model Performance

models = pd.DataFrame({
    'Models': ['Random Forest Classifier', 'Decision Tree Classifier', 'Support Vector Machine',
               'K-Near Neighbors', 'Logistic Model', 'Gaussian NB', 'XGBoost', 'Gradient Boosting'],
    'Score': [RFCCV, DTREECV, SVCCV, KNNCV, LOGCV, GAUSIAN, XGB, GBKCV]})

models.sort_values(by='Score', ascending=False)
```

Out[70]:

	Models	Score
7	Gradient Boosting	0.906217
6	XGBoost	0.905571
0	Random Forest Classifier	0.900093
3	K-Near Neighbors	0.894750
4	Logistic Model	0.892164
1	Decision Tree Classifier	0.874469
5	Gaussian NB	0.843503
2	Support Vector Machine	0.834383

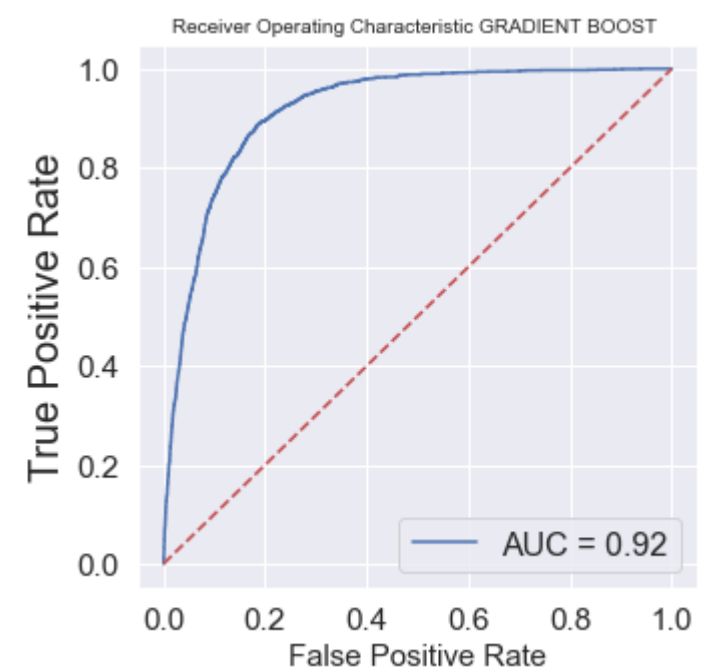
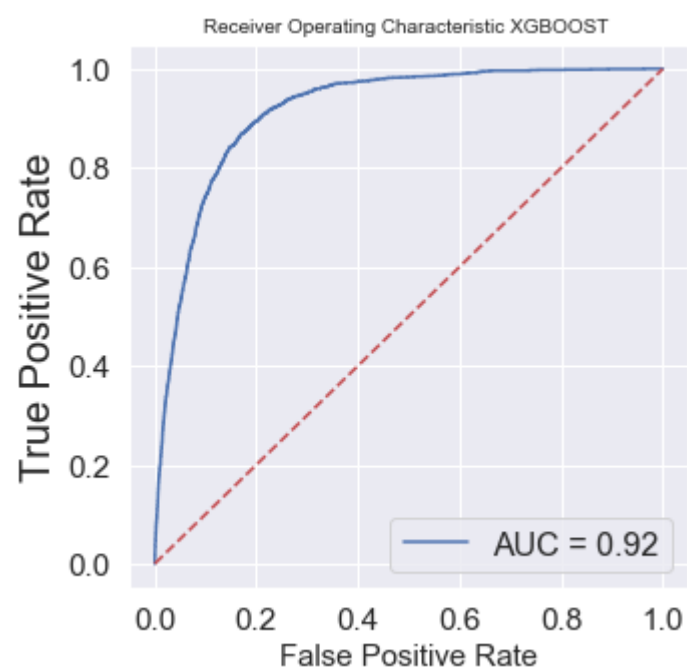
```
In [71]: # XGBOOST ROC/ AUC , BEST MODEL
from sklearn import metrics
fig, (ax, ax1) = plt.subplots(nrows = 1, ncols = 2, figsize = (15,5))
probs = xgb.predict_proba(X_test)
preds = probs[:,1]
fprxgb, tprxgb, thresholdxgb = metrics.roc_curve(y_test, preds)
roc_aucxgb = metrics.auc(fprxgb, tprxgb)

ax.plot(fprxgb, tprxgb, 'b', label = 'AUC = %0.2f' % roc_aucxgb)
ax.plot([0, 1], [0, 1], 'r--')
ax.set_title('Receiver Operating Characteristic XGBOOST ', fontsize=10)
ax.set_ylabel('True Positive Rate', fontsize=20)
ax.set_xlabel('False Positive Rate', fontsize=15)
ax.legend(loc = 'lower right', prop={'size': 16})

#Gradient
probs = gbk.predict_proba(X_test)
preds = probs[:,1]
fprgbk, tprgbk, thresholdgbk = metrics.roc_curve(y_test, preds)
roc_aucgbk = metrics.auc(fprgbk, tprgbk)

ax1.plot(fprgbk, tprgbk, 'b', label = 'AUC = %0.2f' % roc_aucgbk)
ax1.plot([0, 1], [0, 1], 'r--')
ax1.set_title('Receiver Operating Characteristic GRADIENT BOOST ', fontsize=10)
ax1.set_ylabel('True Positive Rate', fontsize=20)
ax1.set_xlabel('False Positive Rate', fontsize=15)
ax1.legend(loc = 'lower right', prop={'size': 16})

plt.subplots_adjust(wspace=1)
```



```
In [72]: #fig, (ax1, ax2, ax3, ax4, ax5) = plt.subplots(nrows = 2, ncols = 3, figsize = (15, 4))
fig, ax_arr = plt.subplots(nrows = 2, ncols = 3, figsize = (20,15))
```

```
#LOGMODEL
probs = logmodel.predict_proba(X_test)
preds = probs[:,1]
fprlog, tprlog, thresholdlog = metrics.roc_curve(y_test, preds)
roc_auclog = metrics.auc(fprlog, tprlog)

ax_arr[0,0].plot(fprlog, tprlog, 'b', label = 'AUC = %0.2f' % roc_auclog)
ax_arr[0,0].plot([0, 1], [0, 1], 'r--')
ax_arr[0,0].set_title('Receiver Operating Characteristic Logistic ', fontsize=20)
ax_arr[0,0].set_ylabel('True Positive Rate', fontsize=20)
ax_arr[0,0].set_xlabel('False Positive Rate', fontsize=15)
ax_arr[0,0].legend(loc = 'lower right', prop={'size': 16})

#RANDOM FOREST -----
probs = classifier.predict_proba(X_test)
preds = probs[:,1]
fprRFC, tprRFC, thresholdRFC = metrics.roc_curve(y_test, preds)
roc_aucRFC = metrics.auc(fprRFC, tprRFC)

ax_arr[0,1].plot(fprRFC, tprRFC, 'b', label = 'AUC = %0.2f' % roc_aucRFC)
ax_arr[0,1].plot([0, 1], [0, 1], 'r--')
ax_arr[0,1].set_title('Receiver Operating Characteristic Random Forest ', fontsize=20)
ax_arr[0,1].set_ylabel('True Positive Rate', fontsize=20)
ax_arr[0,1].set_xlabel('False Positive Rate', fontsize=15)
ax_arr[0,1].legend(loc = 'lower right', prop={'size': 16})

#KNN-----
probs = classifier2.predict_proba(X_test)
preds = probs[:,1]
fprKNN, tprKNN, thresholdKNN = metrics.roc_curve(y_test, preds)
roc_aucKNN = metrics.auc(fprKNN, tprKNN)

ax_arr[0,2].plot(fprKNN, tprKNN, 'b', label = 'AUC = %0.2f' % roc_aucKNN)
ax_arr[0,2].plot([0, 1], [0, 1], 'r--')
ax_arr[0,2].set_title('Receiver Operating Characteristic KNN ', fontsize=20)
ax_arr[0,2].set_ylabel('True Positive Rate', fontsize=20)
ax_arr[0,2].set_xlabel('False Positive Rate', fontsize=15)
ax_arr[0,2].legend(loc = 'lower right', prop={'size': 16})

#DECISION TREE -----
probs = classifier1.predict_proba(X_test)
preds = probs[:,1]
fprDT, tprDT, thresholdDT = metrics.roc_curve(y_test, preds)
roc_aucDT = metrics.auc(fprDT, tprDT)

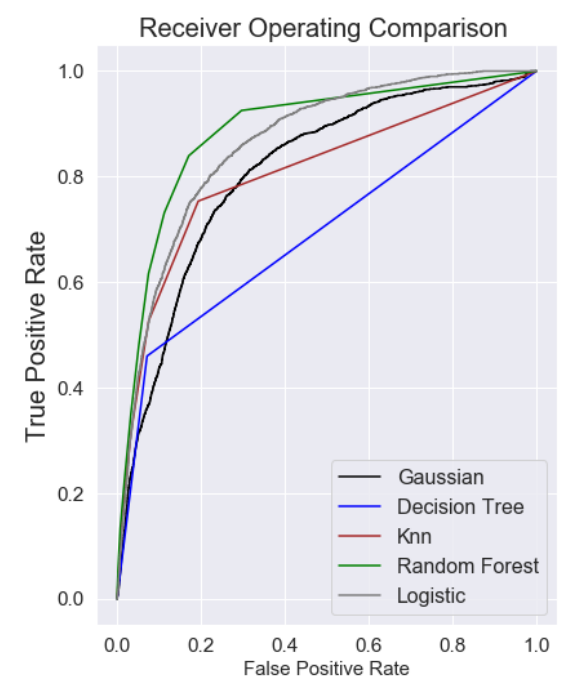





ax_arr[1,0].plot(fprDT, tprDT, 'b', label = 'AUC = %0.2f' % roc_aucDT)
ax_arr[1,0].plot([0, 1], [0, 1], 'r--')
ax_arr[1,0].set_title('Receiver Operating Characteristic Decision Tree ', fontsize=20)
ax_arr[1,0].set_ylabel('True Positive Rate', fontsize=20)
ax_arr[1,0].set_xlabel('False Positive Rate', fontsize=15)
ax_arr[1,0].legend(loc = 'lower right', prop={'size': 16})

#GAUSSIAN -----
probs = gaussianNB.predict_proba(X_test)
preds = probs[:,1]
fprGau, tprGau, thresholdGau = metrics.roc_curve(y_test, preds)
roc_aucGau = metrics.auc(fprGau, tprGau)

ax_arr[1,1].plot(fprGau, tprGau, 'b', label = 'AUC = %0.2f' % roc_aucGau)
ax_arr[1,1].plot([0, 1], [0, 1], 'r--')
ax_arr[1,1].set_title('Receiver Operating Characteristic Gaussian ', fontsize=20)
ax_arr[1,1].set_ylabel('True Positive Rate', fontsize=20)
ax_arr[1,1].set_xlabel('False Positive Rate', fontsize=15)
ax_arr[1,1].legend(loc = 'lower right', prop={'size': 16})

#ALL PLOTS -----
ax_arr[1,2].plot(fprGau, tprGau, 'b', label = 'Gaussian', color='black')
ax_arr[1,2].plot(fprDT, tprDT, 'b', label = 'Decision Tree', color='blue')
ax_arr[1,2].plot(fprKNN, tprKNN, 'b', label = 'Knn', color='brown')
ax_arr[1,2].plot(fprRFC, tprRFC, 'b', label = 'Random Forest', color='green')
ax_arr[1,2].plot(fprlog, tprlog, 'b', label = 'Logistic', color='grey')
ax_arr[1,2].set_title('Receiver Operating Comparison ', fontsize=20)
ax_arr[1,2].set_ylabel('True Positive Rate', fontsize=20)
ax_arr[1,2].set_xlabel('False Positive Rate', fontsize=15)
ax_arr[1,2].legend(loc = 'lower right', prop={'size': 16})

plt.subplots_adjust(wspace=0.2)
plt.tight_layout()
```



In [ ]: